DS-2002: Data Science Systems

A Survey of Data Management Systems

Prof. Jason Williamson

UNIVERSITY *of* VIRGINIA

Class Mechanics

A Survey of Data Management Systems

Who am I?

Stuff about me, but mostly it just impresses my mom

- In addition to UVA, I'm the Executive Director for Industry Solutions

- Prior to that, I ran Oracle for Startups and Oracle for Research for 7 years under Larry Ellison

- Big Company Alum from: Amazon, GE Capital, Capital One

- Small Stuff: Startup Founder, Faith Based NGO, boutique consulting

- Academic: Current BoV at NCCU, Former Prof @McIntire & SEBTS, and UNC-CH Adam's Fellowship Entrepreneurship Mentor

- US Marine and dual citizen of Raleigh, NC and C'Ville VA

- 4 Kids and 1 wife (Anna CAS '23, Patrick CAS '24, Lily (JMU), Jack and Susan COMM '95)

- Can't get enough basketball

- Wrote some books on data, getting a job in big data, stuff like that…again mostly just my mom buys them

- **I love helping students think about and navigate jobs, careers and life --> I've Hired 4 people from this program**

Class Mechanics: Course Content

Lectures:
- In-Person
- First Class of the Week (Tuesday)

Discussion:
- Tuesdays as a part of Lecture Day
- Discord

Readings / Videos:
- Articles and Posts
- Supporting tools and videos
- ~Hours per Week Prep for Tuesdays Class

Labs:
- Hands-On, Virtual & Thursday.
- I may Provide a Zoom invite

Class Mechanics: Course Content

- **Labs**
  - Frequent Hands-On Labs
  - Microsoft Azure Labs (https://labs.azure.com)
  - Begin on Thursdays and Complete by Following Week

Class Mechanics: Grading

- **Data Projects**
  - Concrete Examples of Implementation (Two in total)
  - Released long before they're due
  - Detailed Instructions will be in Canvas (GitHub)
  - No Time Limits!

Class Mechanics: Grading

Component	Weight	Frequency
Student Survey		First Week
Lectures / Readings / Material		Weekly
Engaged Discussion	5%	Weekly
Quiz	15%	3 or 4
Labs	30%	Weekly
Data Projects	50%	2 projects (25% each)

Class Mechanics: Things You Need to Know

How this course will be conducted: Class Location, Course Materials, Communications

- This class will be taught IN-PERSON
- Wed/Friday we may be on Zoom for Hands-on Labs
    - Sometimes we'll do them in groups
    - …And sometimes you'll do them alone
- The schedule is prone to updates: I'll be sure to keep everyone informed!
- Everything will be updated on GitHub:

- **My Repo will be here --> https://github.com/jasonwnc/ds2002F23**

- Quizzes, Projects and Grades will be released on Canvas
- You can DM me on Discord with Questions or to Set-up Office Hours

Learning Objectives

Develop Robust Facility for Handing Data

Acquire a Strong Understanding of SQL and NoSQL Databases

Understand Systems Options in the Public Cloud

Gain Fluency in Tools, Processes, and Services …and How to Choose Among Them

Data Retrieval

Data Shipping

Data Schemas

Data Ingestion

Data Processing

Data Normalization

Data Analysis

In the Beginning… There Was Mainframe

- Data stored in ISAM files (indexed sequential access method)
  - Flat files having fixed-length fields
- Centralized computing and storage resources
  - "Dumb" terminal clients
- Only affordable to large, wealthy enterprises:
  - Governments
  - Large Banks

Then Came... Client/Server Networks

- Data stored in Relational Database Management Systems (RDBMS):
  - Oracle
  - SQL Server
- Enterprise Data Centers:
  - Company owned & managed
  - Based on commodity hardware
- Some computation and storage resources shared by clients

Which Gave Us… SMP Servers

Relational Database Management running on **Symmetric Multi-Processing Servers**

- Dedicated Database Servers:
  - RDBMS (Relational Database Management Server)
  - Oracle, SQL Server, DB2, etc.
- OLTP (Online Transaction Processing)
  - Characterized by a large volume of transactions each of which affect a small number of rows
  - Online Sales, Bank Deposits & Transfers
  - Highly Normalized Database Schema
- OLAP (Online Analytical Processing)
  - Characterized by a small volume of read transactions each of which affect a large number of rows
  - Periodic Post-hoc Analysis *(What Happened?)*
  - De-Normalized Multi-Dimensional Data Warehouse Schema

But Then… An Explosion of Data – "BIG Data"
A Rapid, Exponential Proliferation of New Devices: The Internet of Things (IoT)

Volume (Size)
- Explosion in social media, mobile apps, digital sensors, RFID, GPS, and more have caused exponential data growth.

Velocity (Speed)
- Sources like social networking and sensor signals create data at a tremendous rate; making it a challenge to capture, store, and analyze that data in a timely or economical manner.

Variety (Structure)
- Traditionally BI has sourced structured data, but now insight must be extracted from unstructured or poly-schematic data like large text blobs, digital media, sensor data, etc.

Veracity (Quality)
- The anonymity of the WWW, incredible sources like social networking and duplicate systems bring into question the authenticity of the information being generated and collected.

But… How Can We Manage All That Data: Scalability

The Advent of Big Data Drove a Need to Increase Scale to the Petabyte Level: $$$$$

SMP: Symmetric Multi-Processing

Scale Up (Limited)

This

Became

MPP: Massively Parallel Processing

Scale Out (Practically Unlimited)

And Today… Cloud-hosted Databases & Services

Re-centralization of Servers ▪ Based on Commodity HW ▪ Resources Shared by Client

Essential Characteristics

- **On-Demand Self-Service**
- **Broad Network Access**
- **Resource Pooling**
- **Rapid Elasticity**
- **Measured Service**

Service Models

- **Infrastructure as a Service (IaaS)**
- **Platform as a Service (PaaS)**
- **Software as a Service (SaaS)**

Deployment Models

- **Public Cloud**
- **Private Cloud**
- **Community Cloud**
- **Hybrid Cloud**

As defined by NIST (National Institute of Standards and Technology)

Service Models: On-Premises vs Cloud-Hosted

Managed by Owner	Managed by Service Provider

On Premises	IaaS	PaaS	SaaS
Applications	Applications	Applications	Applications
Data	Data	Data	Data
Runtime	Runtime	Runtime	Runtime
Middleware	Middleware	Middleware	Middleware
OS	OS	OS	OS
Virtualization	Virtualization	Virtualization	Virtualization
Servers	Servers	Servers	Servers
Storage	Storage	Storage	Storage
Networking	Networking	Networking	Networking

Data Management… In the Cloud
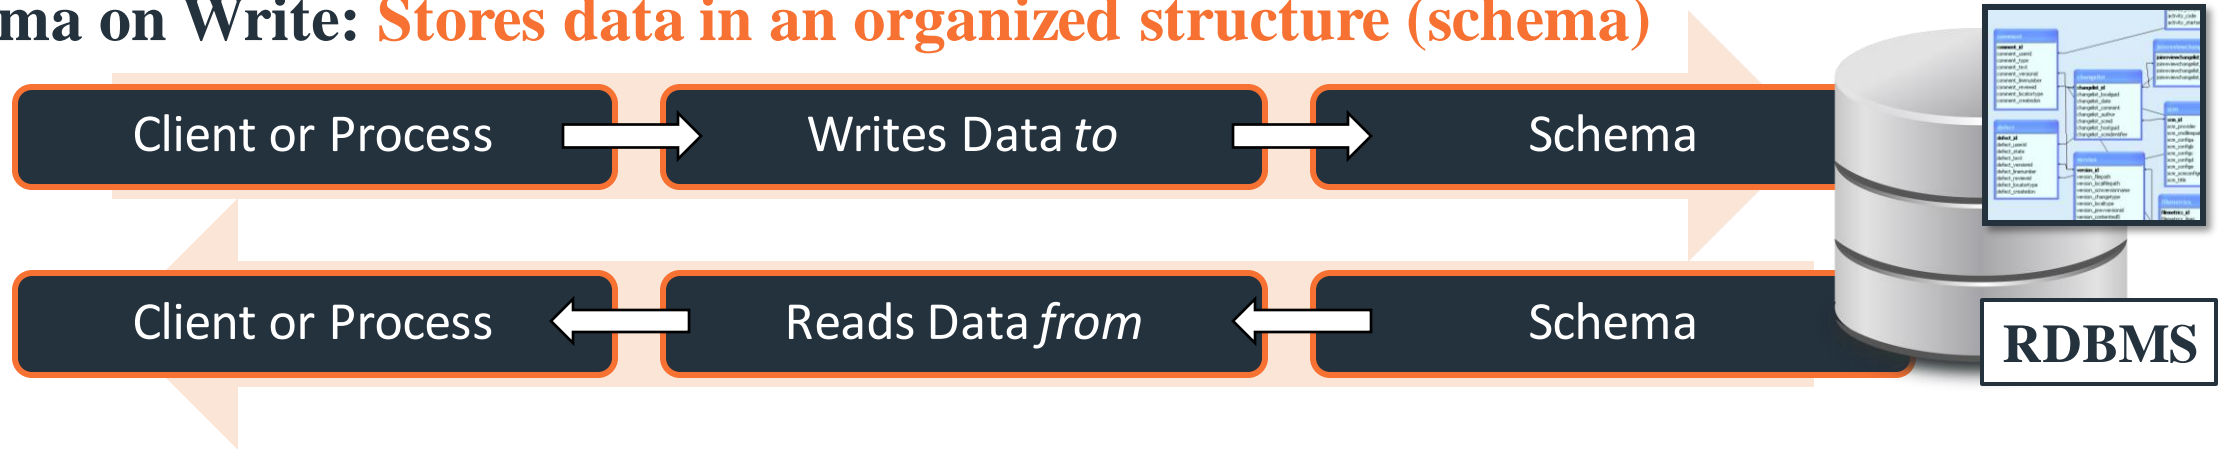Economies of Scale Afforded by the Cloud Have Enabled Massive Storage & Compute

- File-based Storage: Blobs, Queues, Tables, Data Lakes
- Relational Database Management Systems:
  - Online Transaction Processing (OLTP) systems:
    - **IaaS**: Cloud-hosted VMs Running RDBMS Software & Databases
    - **PaaS:** (Database as a Service) Single Databases, Multi-Database Pools, Managed Instance
  - Online Analytical Processing (OLAP) systems:
    - **PaaS:** Massively Parallel Processing (MPP) Data Warehouse
- NoSQL Database Management Systems:
  - Semi-Structured or Poly-Schematic data
  - Massively Parallel Processing
    - Data partitioned and replicated across many server machines (nodes)
    - Data typically stored in file-based format (e.g., JavaScript Object Notation (JSON)
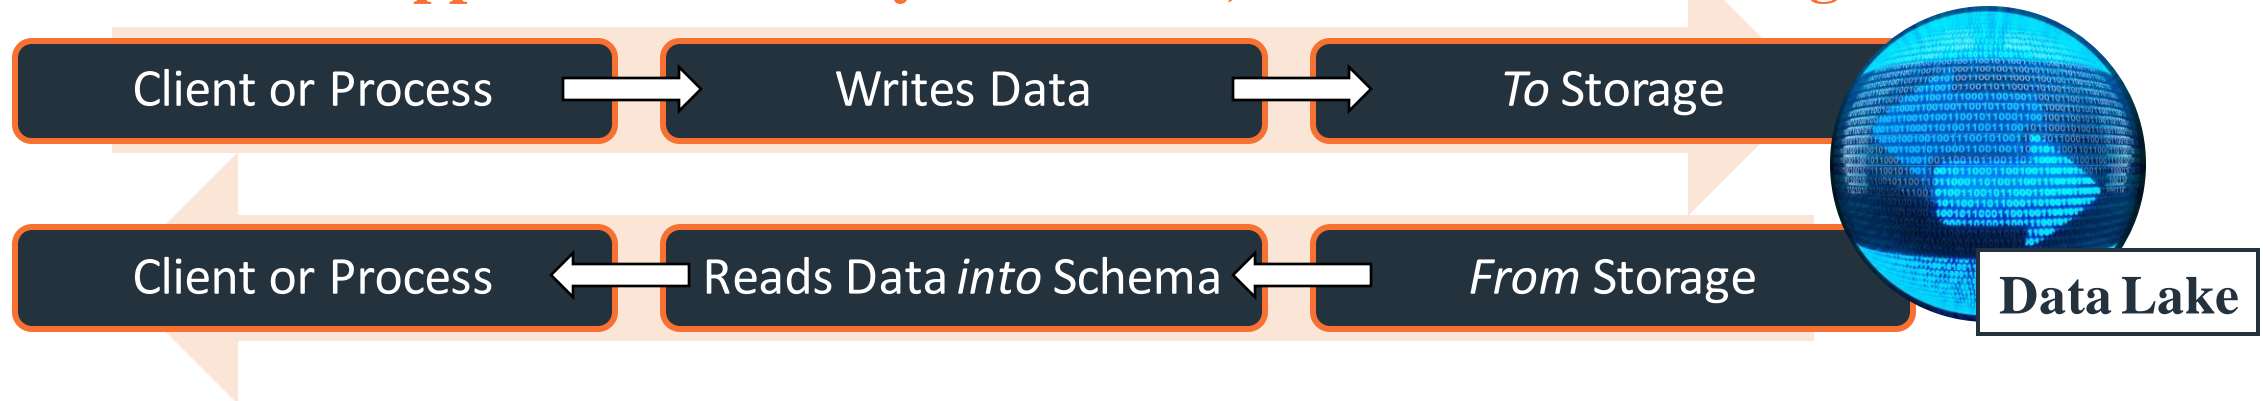- Data Lakehouse: Data Warehouse Schema over File-based Storage (Data Lake)

Paradigms: Data Storage and Retrieval
Schema on Write versus Schema on Read

Schema on Write: Stores data in an organized structure (schema)

| Client or Process | → | Writes Data *to* | → | Schema |

| Client or Process | ← | Reads Data *from* | ← | Schema |

RDBMS

Schema on Read: Applies schema only when read, data stored in its original format

| Client or Process | → | Writes Data | → | *To* Storage |

| Client or Process | ← | Reads Data *into* Schema | ← | *From* Storage |

Data Lake

NoSQL: Not Only SQL

Key-value stores
- The simplest NoSQL database; based on "dictionaries" or "maps".
- Items are stored in associative arrays; pairing a name (or "key"), with a value.
- **Riak, FoundationDB, and Redis**

Column stores
- Combines a key, value and timestamp for each item.
- Optimized for large datasets by storing columns of data together, rather than in rows.
- **Cassandra, BigTable and HBase**

Document databases
- Pairs keys with complex data structures (documents) using XML, JSON or BSON.
- Documents may contain key-value pairs, key-array pairs, and nested documents.
- **MongoDB, MarkLogic, and Apache CouchDB**

Graph stores
- Stores interrelated networks of data such as social connections, or network topologies.
- Optimized for interconnected data elements with an undetermined number of relations.
- **AllegroGraph, Neo4J and HyperGraphDB**.
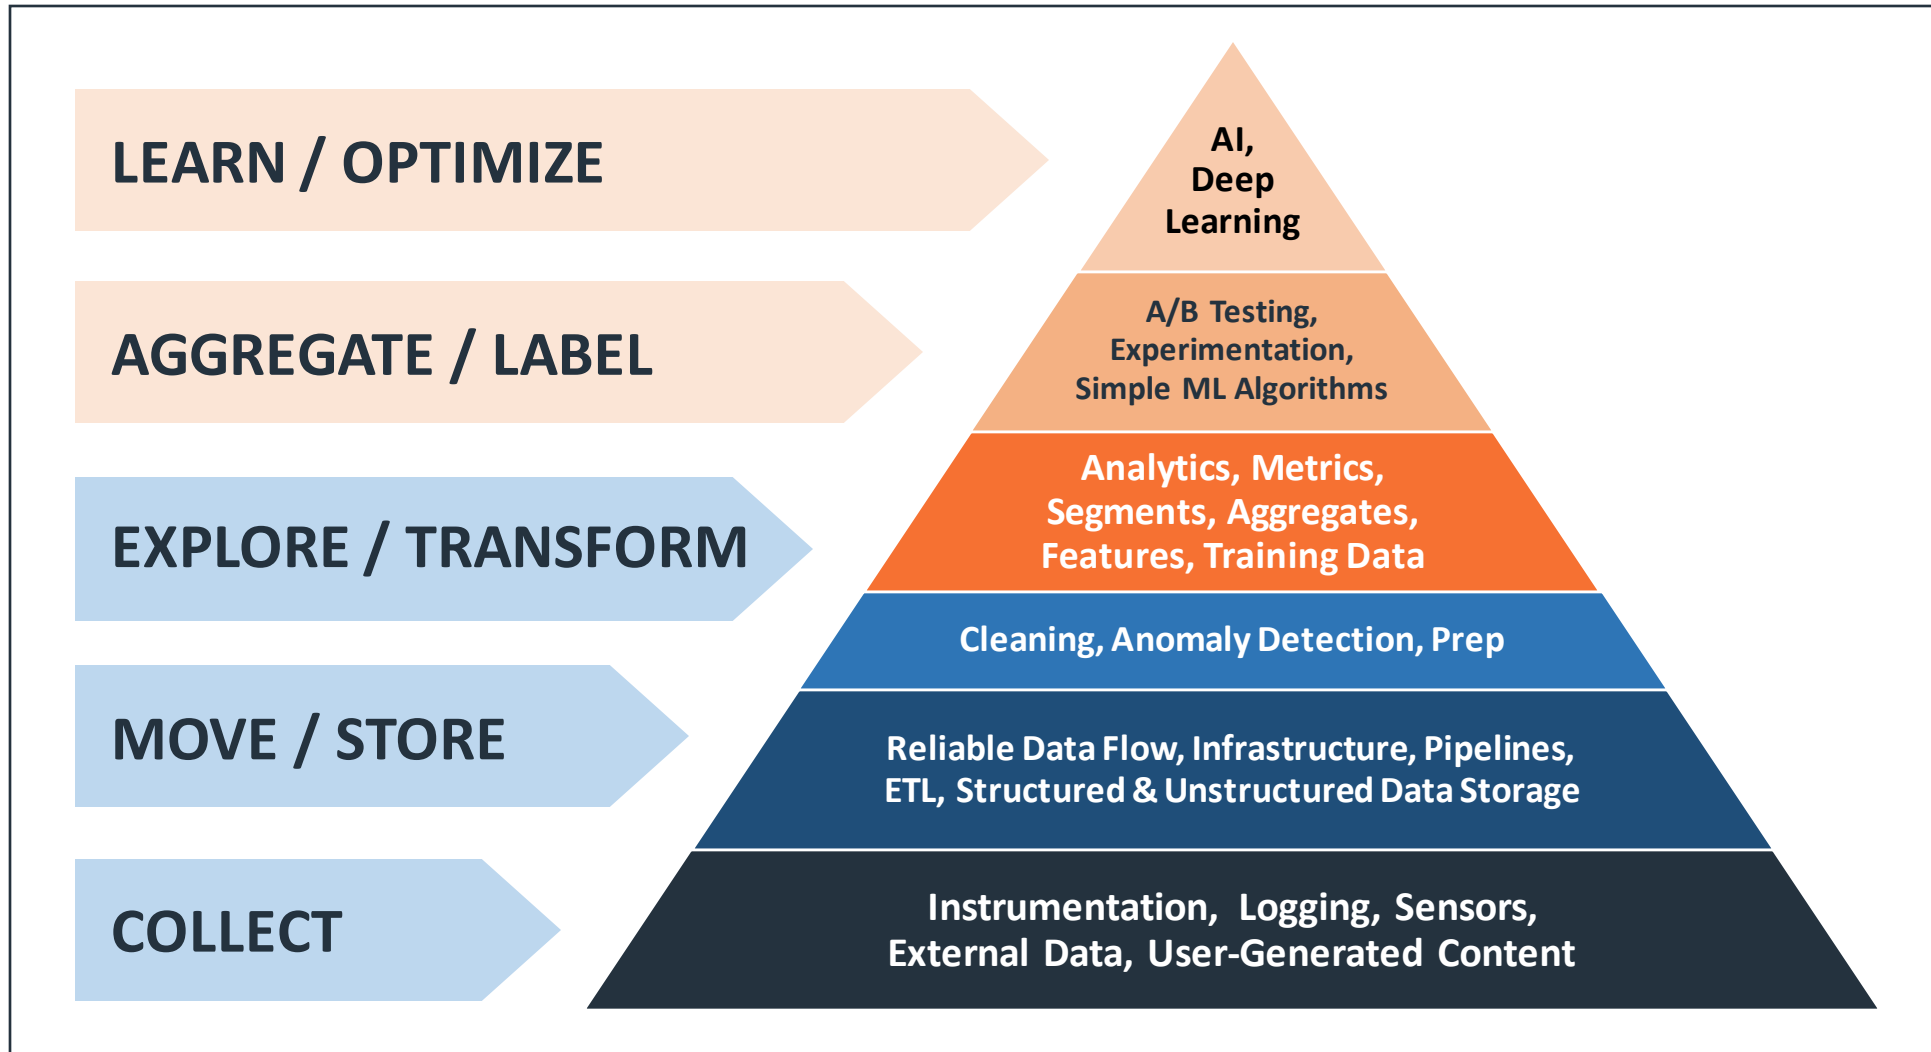
SQL versus NoSQL Databases

	SQL Databases	NoSQL Databases
Types	One: Relational Database	Many: Key-value, document, column, and graph databases
Data Storage	Records are stored as rows in tables that represent entities. Entity relationships are modeled by joining tables.	Records may be stored as "documents" using XML or JSON. Entity relationships are modeled by nesting them hierarchically.
Schemas	Static: Structure & data types are fixed at design time. Adding new elements requires schema design changes.	Dynamic: New elements can be added at runtime. Poly-schematic: Dissimilar data can be stored together as necessary.
Scaling	Vertically: A single server must be made increasingly powerful to cope with increasing demand.	Horizontally: New commodity servers and storage are added to an array. Data is automatically distributed across all servers.
Transactions	Full transactional consistency (ACID)	Single element (document) only.
Manipulation	Using platform-specific languages (SQL)	Using OSS API's, low-level lang., JavaScript
Consistency	Supports strong consistency	Per-product: Most are eventually consistent.

Onward to Data Science

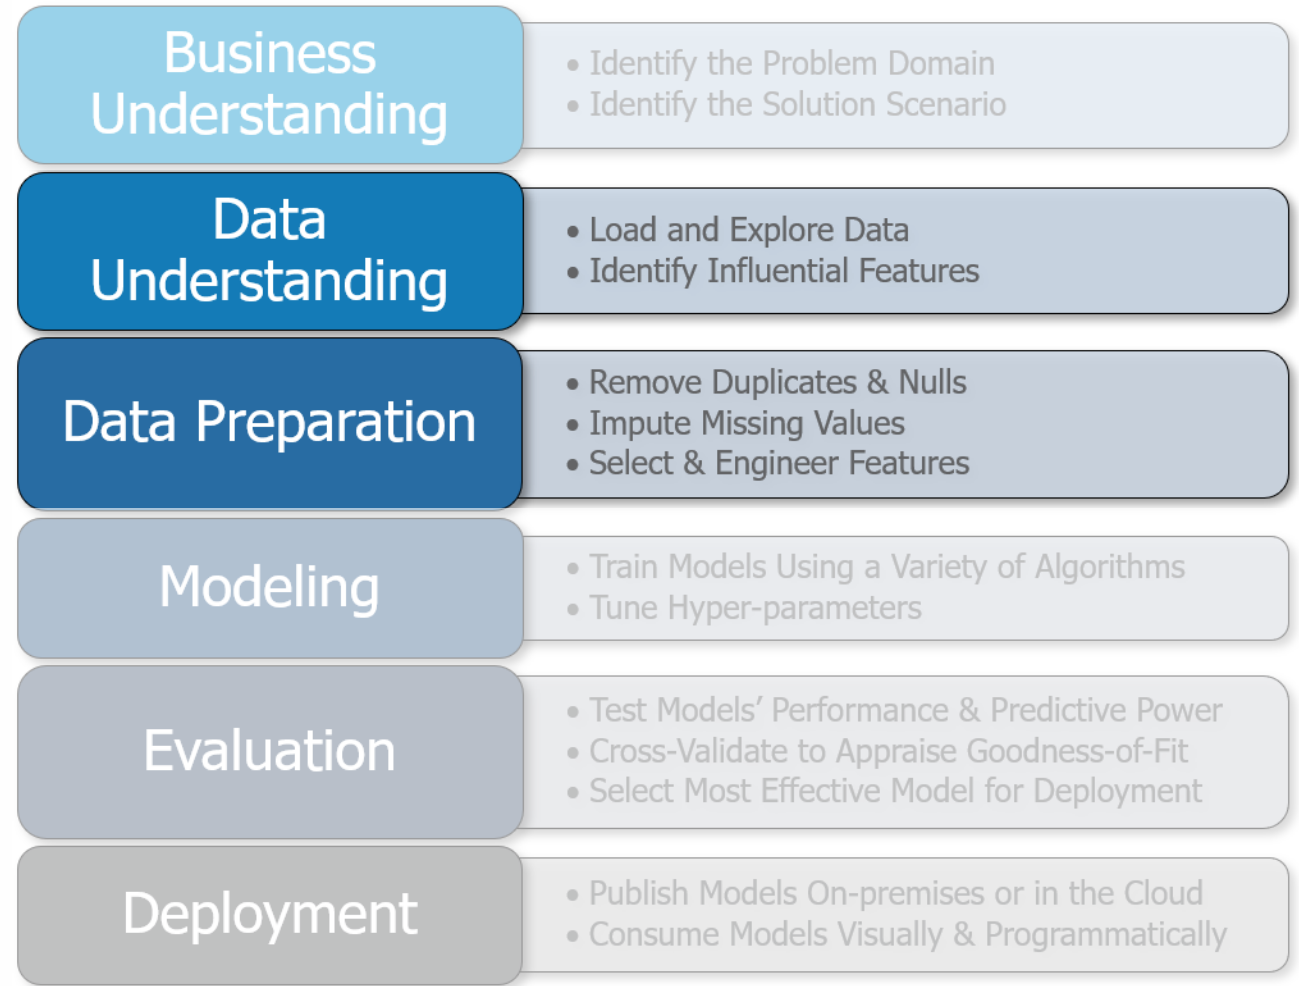A Brief Overview of How Data Systems are Involved

UNIVERSITY*of*VIRGINIA

CRISP-DM: Cross-Industry Standard Process-Data Mining

First Introduced in 1996!

Business Understanding
- Identify the Problem Domain
- Identify the Solution Scenario

Data Understanding
- Load and Explore Data
- Identify Influential Features

Data Preparation
- Remove Duplicates & Nulls
- Impute Missing Values
- Select & Engineer Features

Modeling
- Train Models Using a Variety of Algorithms
- Tune Hyper-parameters

Evaluation
- Test Models' Performance & Predictive Power
- Cross-Validate to Appraise Goodness-of-Fit
- Select Most Effective Model for Deployment

Deployment
- Publish Models On-premises or in the Cloud
- Consume Models Visually & Programmatically

Data Engineering... for Data Science

Frequently, Data Must Be Moved from Sources to a Database and/or Data Lake

Extract

- This is the step where sensors wait for upstream data sources to land. Once available, we transport the data from their source locations to further transformations.
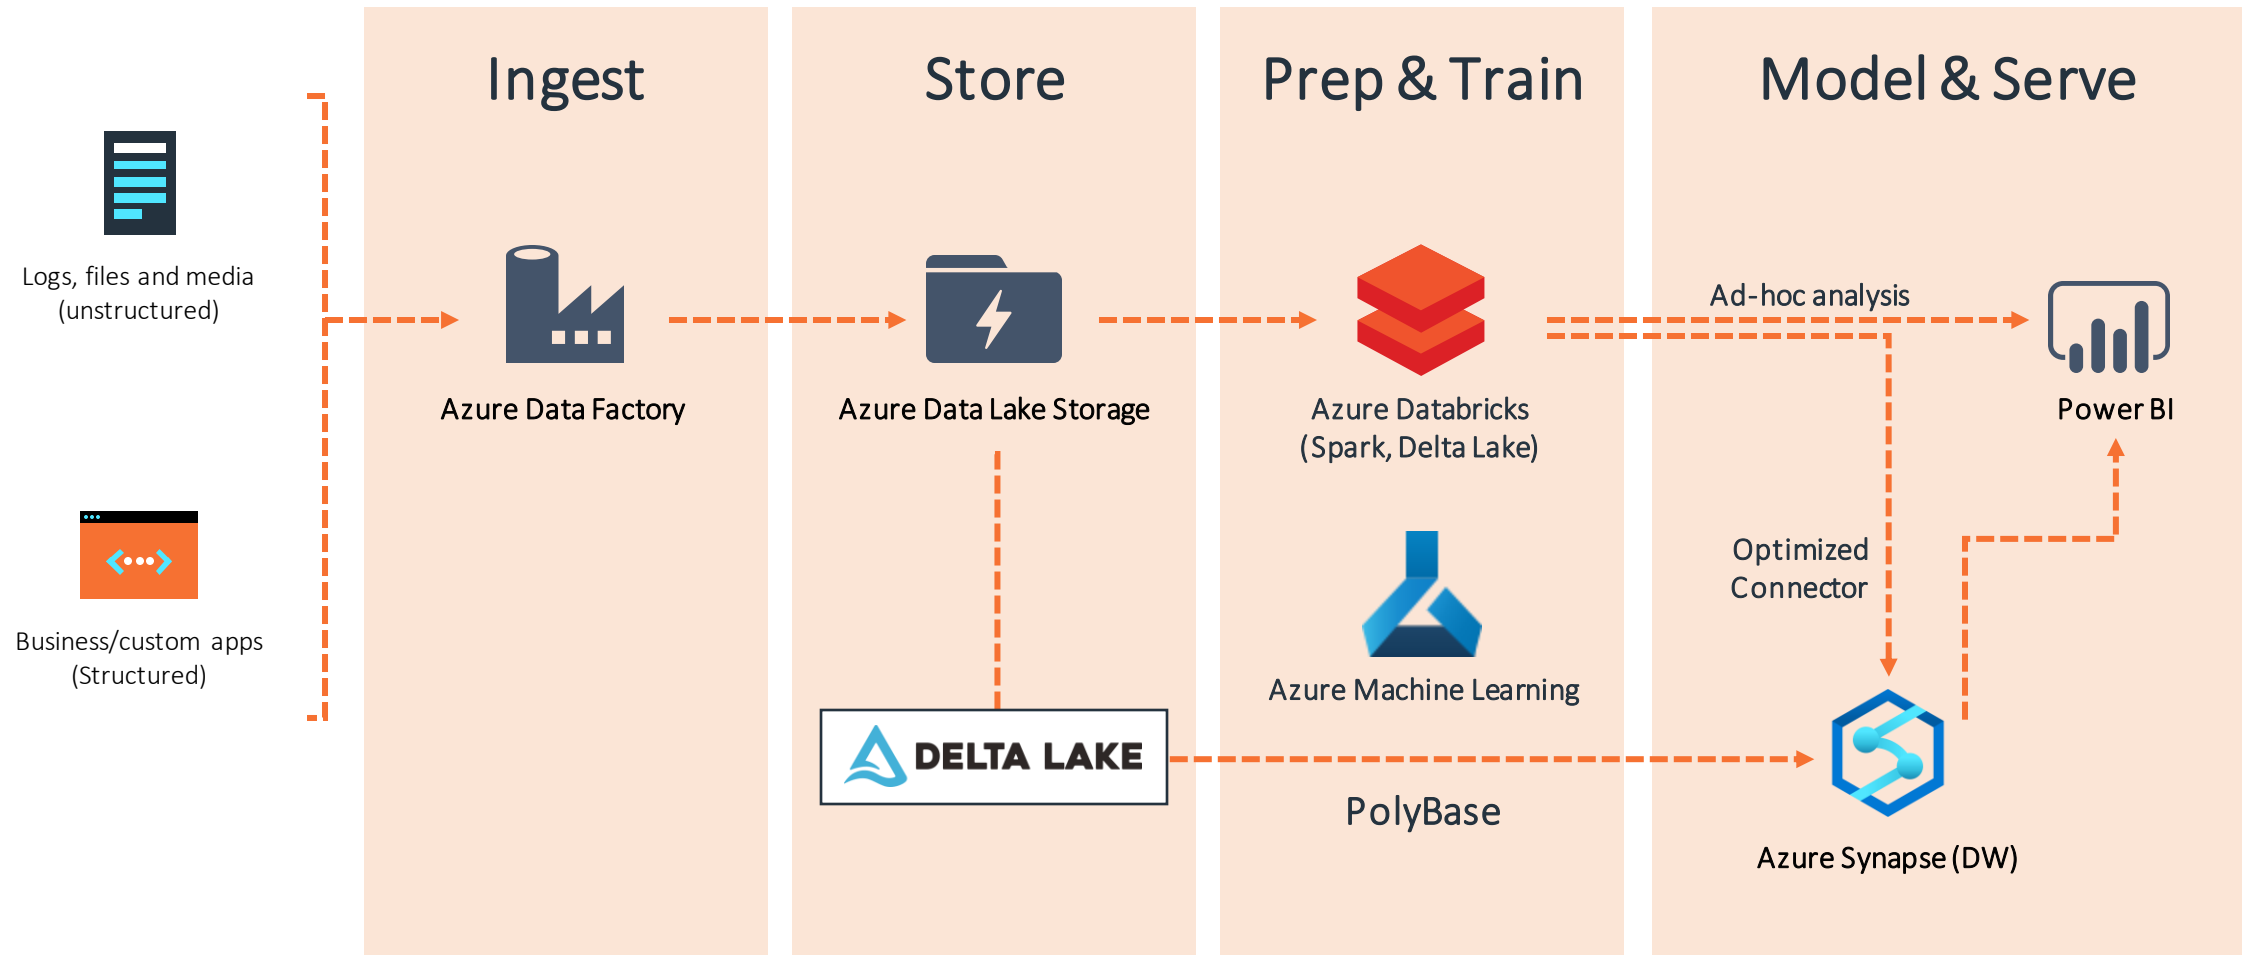
Transform

- The heart of any ETL job: apply business logic, perform actions such as filtering, grouping, and aggregation to translate raw data into analysis-ready datasets.
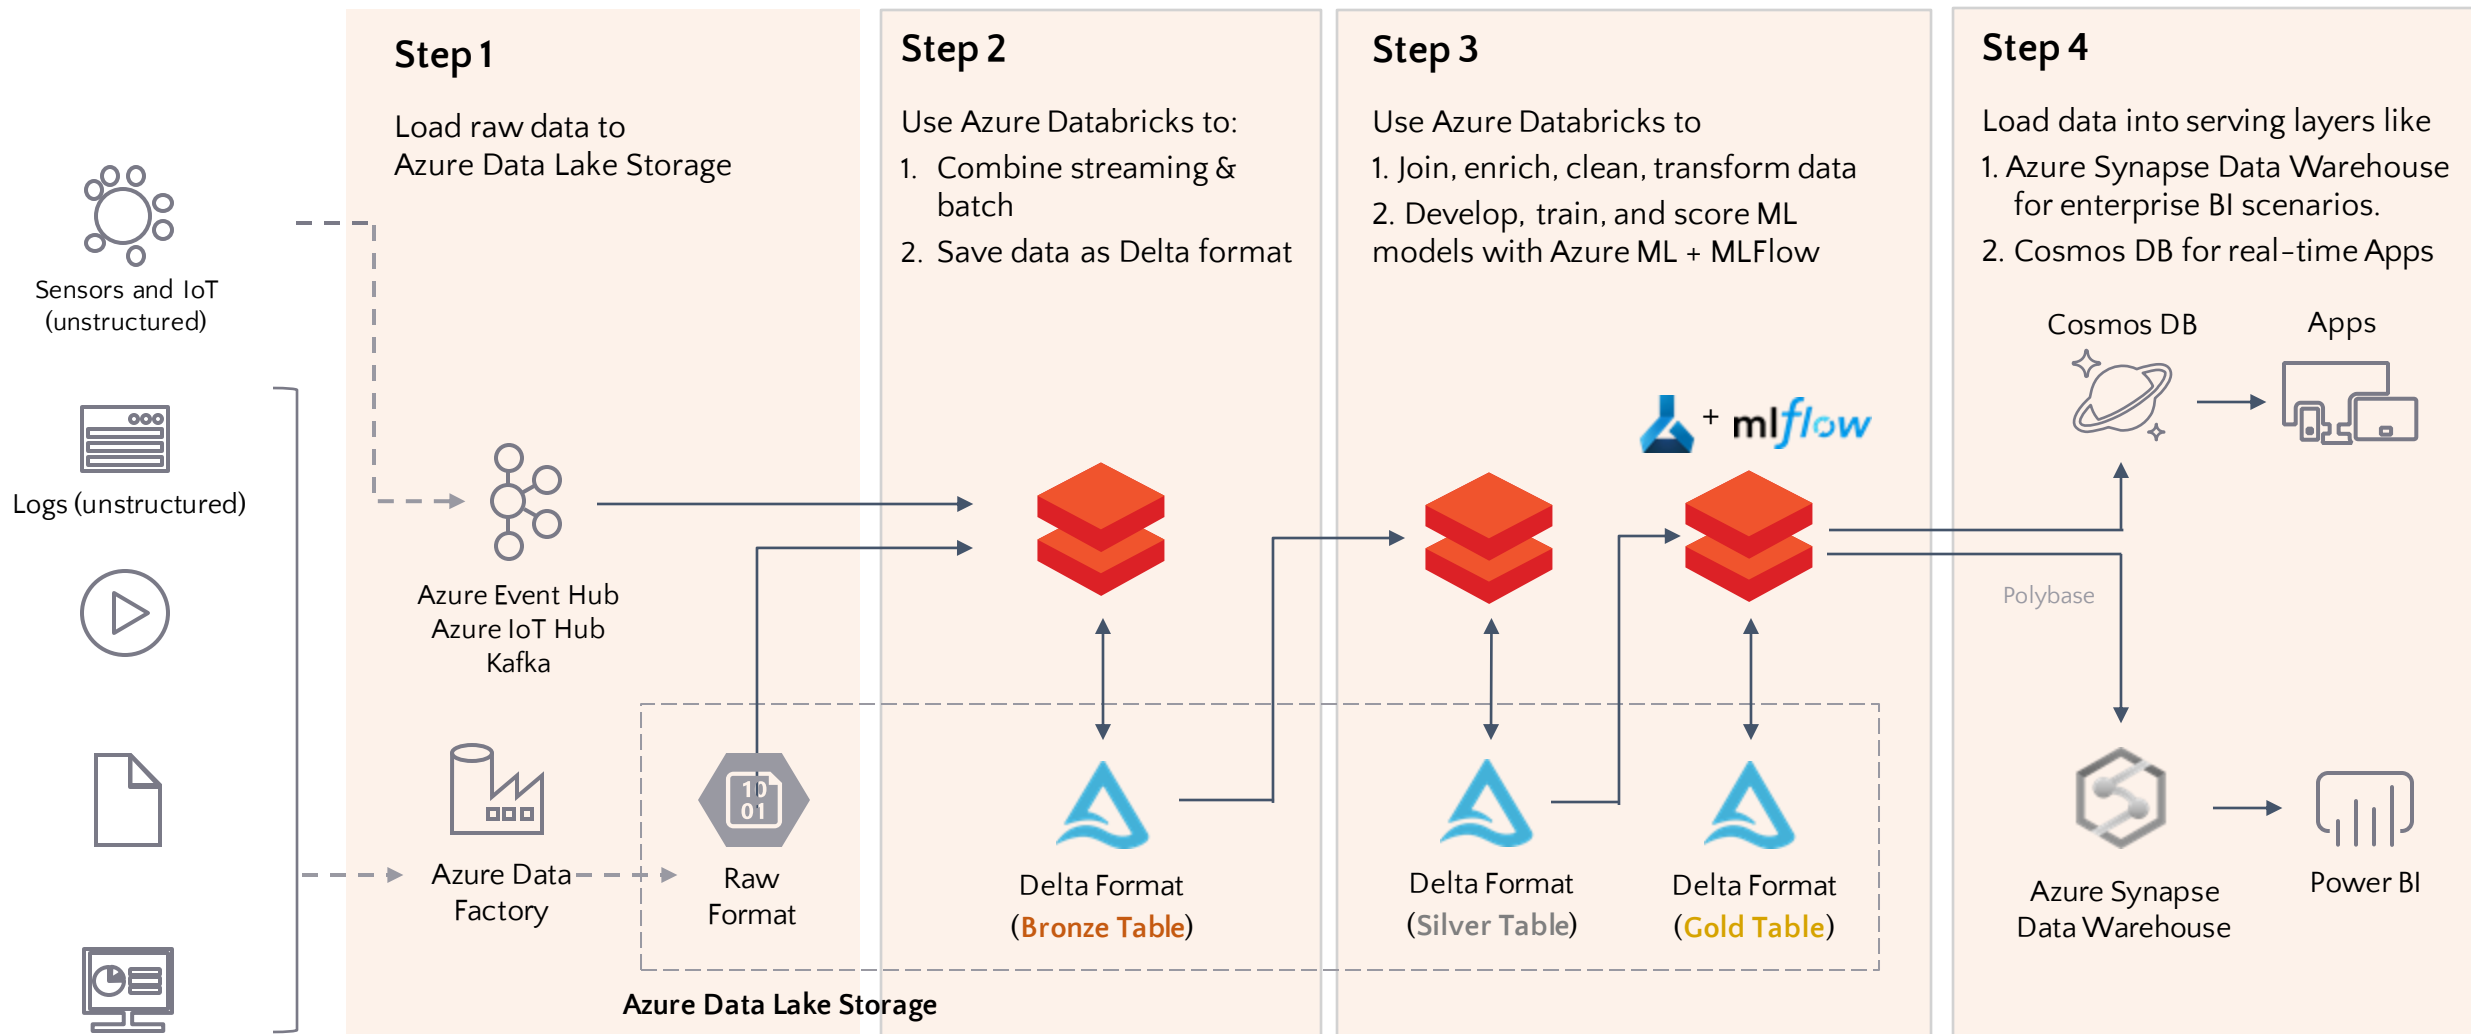
Load

- Load the processed data and transport to a final destination. Can now be consumed directly by end-users or treated as yet another upstream dependency.

Design Pattern: Modern Data Warehousing

Step 1

Load raw data to Azure Data Lake Storage

Step 2

Use Azure Databricks to:
1. Combine streaming & batch
2. Save data as Delta format

Step 3

Use Azure Databricks to
1. Join, enrich, clean, transform data
2. Develop, train, and score ML models with Azure ML + MLFlow

Step 4

Load data into serving layers like
1. Azure Synapse Data Warehouse for enterprise BI scenarios.
2. Cosmos DB for real-time Apps

Sensors and IoT (unstructured)

Logs (unstructured)

Azure Event Hub
Azure IoT Hub
Kafka

Azure Data Factory

Raw Format

Delta Format (**Bronze Table**)

Delta Format (**Silver Table**)

Delta Format (**Gold Table**)

Polybase

Cosmos DB

Apps

Azure Synapse Data Warehouse

Power BI

Azure Data Lake Storage

Q & A

A Survey of Data Management Systems